



# **Lecture 18: Real-Time Process Scheduling**

EECS 388 – Fall 2022

© Prof. Mohammad Alian

Lecture notes are based on slides created by Prof.  
Heechul Yun

# Recap

- Job

- A computation instance

- Task

- A sequence of jobs

- Periodic task model

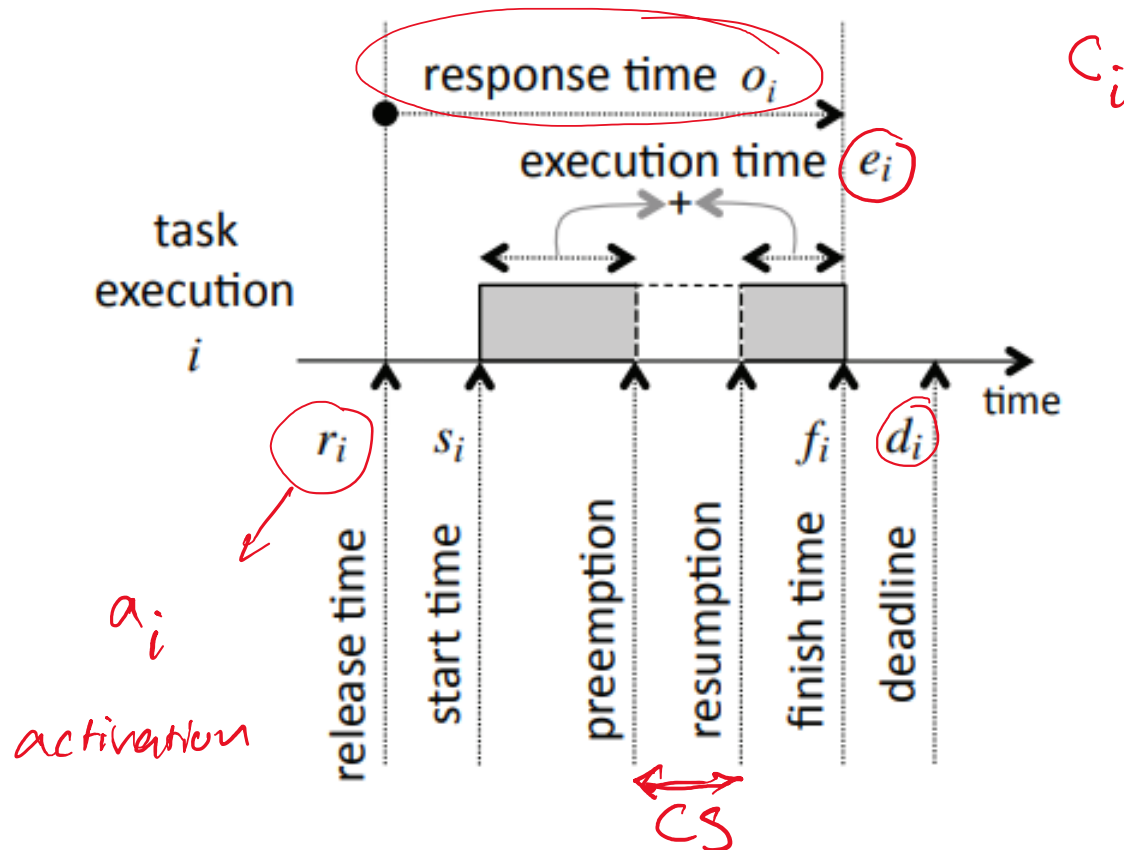
- $t_i = (\underline{C_i}, \underline{T_i})$  or  $(C_i, T_i, D_i)$

Recommended Reading: Chapter 12

[Introduction to Embedded Systems - A Cyber-Physical Systems Approach](#), by E Lee and S.

Seshia.

# Times Associated with a Task Execution – Text Book's Terminology




# Real-Time Scheduling

- Scheduling
  - Pick which task to run next
- Priority-based scheduling
  - Pick the highest priority task among ready tasks

# Priority and Criticality

- Priority
  - priority is the order we execute ready jobs.
- Criticality (Importance)
  - represents the penalty if a task misses a deadline (one of its jobs misses a deadline).
- Question
  - Which task should have higher priority?
  - Task 1: The most important task in the system: if it does not get done, serious consequences will occur
  - Task 2: A mp3 player: if it does not get done in time, the played song will have a glitch
  - If it is feasible, we would like to meet the real-time deadlines of both tasks

# Priority and Criticality




- If priorities are assigned according to importance, you can miss the deadline even if the system is mostly idle
- If  $p_2$  (period of  $t_2$ )  $<$   $C_1$  (execution time of  $t_1$ ),  $t_2$  will miss the deadline

# Utilization

- A task's utilization (of the CPU) is the fraction of time for which it uses the CPU (over a long interval of time).
- A periodic task's utilization  $U_i$  (of CPU) is the ratio between its execution time and period:  $U_i = C_i/p_i$
- Given a set of periodic tasks, the total CPU's utilization is equal to the sum of periodic tasks' utilization:

$$U = \sum_i \frac{C_i}{p_i}$$

single CPU



- What is the limit on  $U$ ?

$$\frac{4}{10}$$

$$U_1 = \frac{3}{10} + U_2 = \frac{1}{10}$$

# Real-Time Scheduling Algorithms

- Fixed (static) -priority scheduling
  - All jobs of a task have the same priority
- • Rate Monotonic (RM)
- Dynamic priority scheduling
  - Different jobs of the same task may have different priorities
  - Earliest Deadline First (EDF)




# Rate Monotonic (RM)

- Fixed-priority scheduling algorithm
- Assigns priorities to tasks on the basis of their periods
- Shorter period = higher priority.

# RM Example

- Case#1

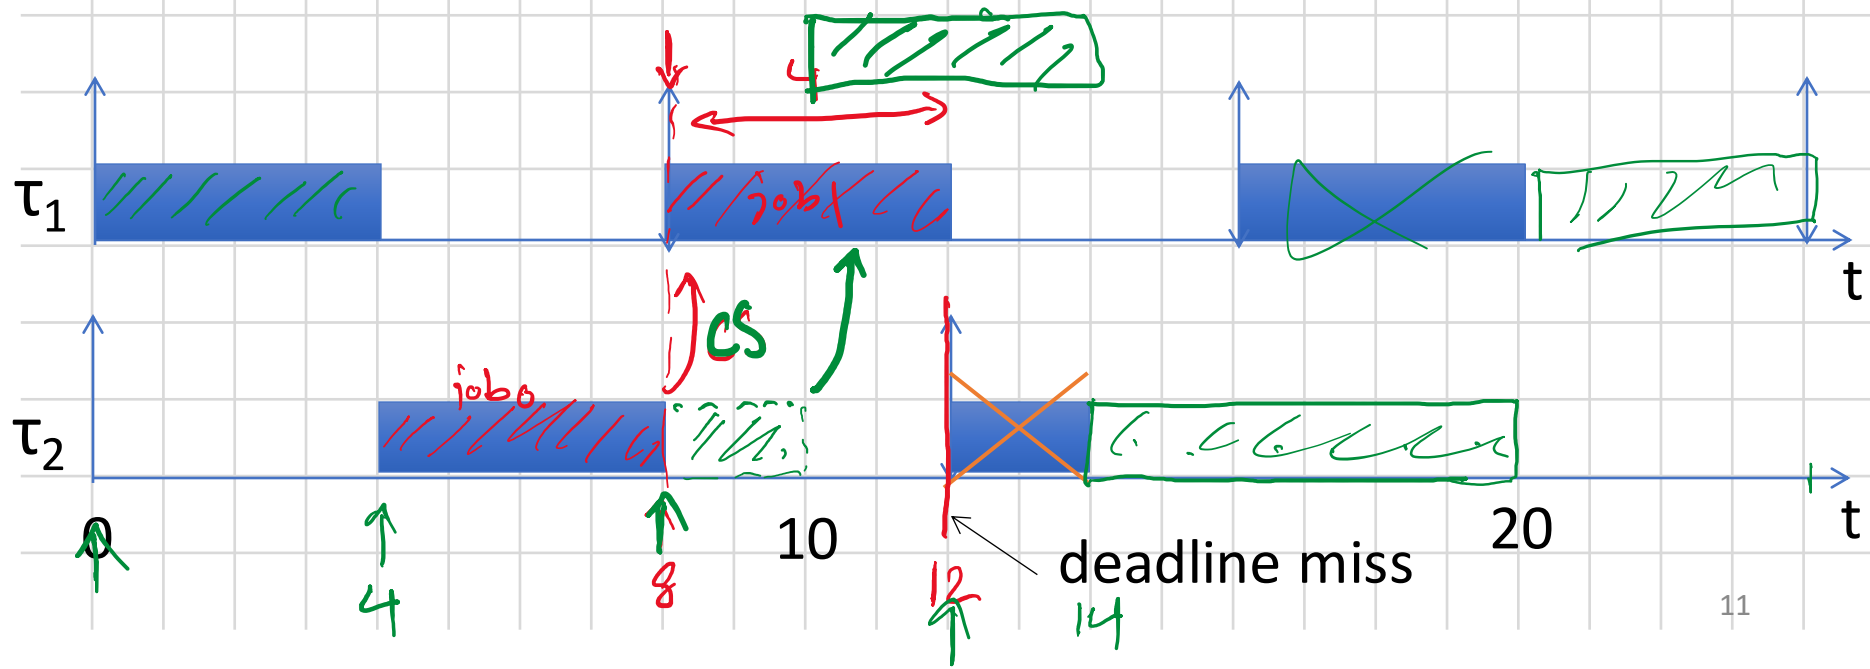
- $\tau_1$  ( $C1 = 4$ ,  $T1 = 8$ ), high prio
- $\tau_2$  ( $C2 = 3$ ,  $T2 = 12$ ), low prio
- Utilization:  $U = \frac{4}{8} + \frac{3}{12} = 0.75$



# RM Example

- Case#2

- $\tau_1$  ( $C1 = 4$ ,  $T1 = 8$ ), high prio
- $\tau_2$  ( $C2 = 6$ ,  $T2 = 12$ ), low prio
- Utilization:  $U = 4/8 + 6/12 = 1$




# Earliest Deadline First (EDF)

- Dynamic-Priority Scheduling Algorithm
- Task priority is inversely proportional to its current absolute deadline
- **Earlier deadline = higher priority**
- Each job of a task has a different deadline, hence a different priority.

# EDF Example

- Case#2:

- $\tau_1$  ( $C1 = 4$ ,  $T1 = 8$ ),  $\tau_2$  ( $C2 = 6$ ,  $T1 = 12$ )
- Utilization:  $U = 4/8 + 6/12 = U_b = 1$



## 1

- 


---

- lowest pri



# RM vs. EDF

- $\tau_1$  ( $C1 = 1, T1 = 4$ ),  $\tau_2$  ( $C2 = 2, T1 = 6$ ),  $\tau_3$  ( $C3 = 3, T3 = 8$ )
- Utilization:  $U = 1/4 + 2/6 + 3/8 = 23/24$
- EDF scheduling?



# Key Results

For periodic tasks with relative deadlines equal to their periods:

- Rate monotonic scheduling is the optimal fixed-priority priority policy
  - No other static priority assignment can do better
  - Yet, it might not achieve 100% CPU utilization
- Earliest deadline first scheduling is the optimal dynamic priority policy
  - EDF can achieve 100% CPU utilization



# RM vs. EDF

- In practice, industrial systems heavily favor RM over EDF. Why?

# RM vs. EDF

- In practice, industrial systems heavily favor RM over EDF. Why?
- RM is easier to implement
  - Task priority never changes.
- RM is more transparent and robust
  - easier to understand what is going on if something goes wrong (ex: overload).
- ➔ • if a task executes for longer than its prescribed worst-case time, higher priority tasks will be left untouched.

# Example

Consider the following real-time taskset:

Task	C	P	U
t1	2	10	0.200
t2	4	15	0.267
t3	10	35	0.286

Note1: C = worst-case execution time; P = period; U = utilization.  
Assume a single core processor.









Handwritten notes: A blue box encloses the P column. Blue circles highlight the C values for t1 and t3. Blue arrows point to the U values, and a blue summation symbol  $\Sigma$  is written below the U column.

1- Compute the total utilization of the taskset.


2- Assign priorities of the tasks following the rate monotonic policy. You can use numbers, e.g. 1 (lowest priority) to 3 (highest).

3- Draw a timeline (from time 0 to 40) of the taskset under the rate monotonic scheduling (RMS).

EDA



RMS



EDF

finish task3

# Acknowledgements

- These slides draw on materials developed by
  - Lui Sha and Marco Caccamo (UIUC)
  - Rodolfo Pellizzoni (U. Waterloo)
  - Edward A. Lee and Prabal Dutta (UCB) for EECS149/249A