Lecture 20: Mutual Exclusion and Scheduling

EECS 388 – Fall 2022

© Prof. Mohammad Alian

Lecture notes are in parts based on slides created by Prof. Heechul Yun

Context

Recommended Reading - Chapter 12.4 and 11.2.3

Lee and Seshia "Introduction to Embedded Systems"

- Add a bit of complexity
  - How resource sharing affect scheduling?

- Mutual Exclusion
- Race Condition, Critical Section, lock

Race Condition

Initial condition: *counter = 5*

- What are the possible outcome?

Race Condition

Initial condition: *counter = 5*

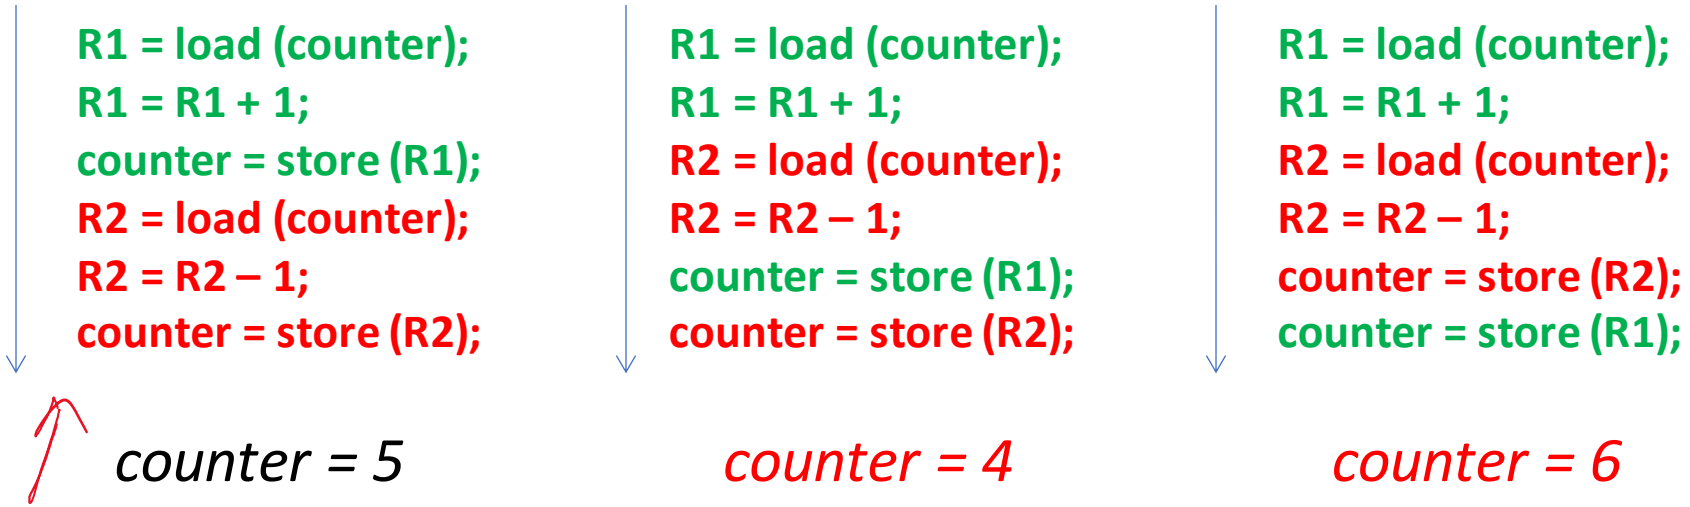

- Why this happens?

Race Condition

- A situation when two or more threads **read and write** shared data at the same time
- Correctness depends on the execution order

Thread 1	Thread 2
R1 = load (**counter**);	R2 = load (**counter**);
R1 = R1 + 1;	R2 = R2 – 1;
counter = store (R1);	**counter** = store (R2);

- How to prevent race conditions?

Critical Section

- Code sections of potential race conditions

Thread 1	Thread 2
Do something	Do something
..	..
R1 = load (**counter**);	R2 = load (**counter**);
R1 = R1 + 1;	R2 = R2 – 1;
counter = store (R1);	**counter** = store (R2);
...	..
Do something	Do something

Critical sections

Mutual Exclusion

- A property that requires only one thread can enter its critical section at a time among multiple concurrent threads
- Lock (mutex) is a mechanism to provide mutual exclusion

Lock

- General solution
  - Protect critical section via a lock
  - Acquire on enter, release on exit

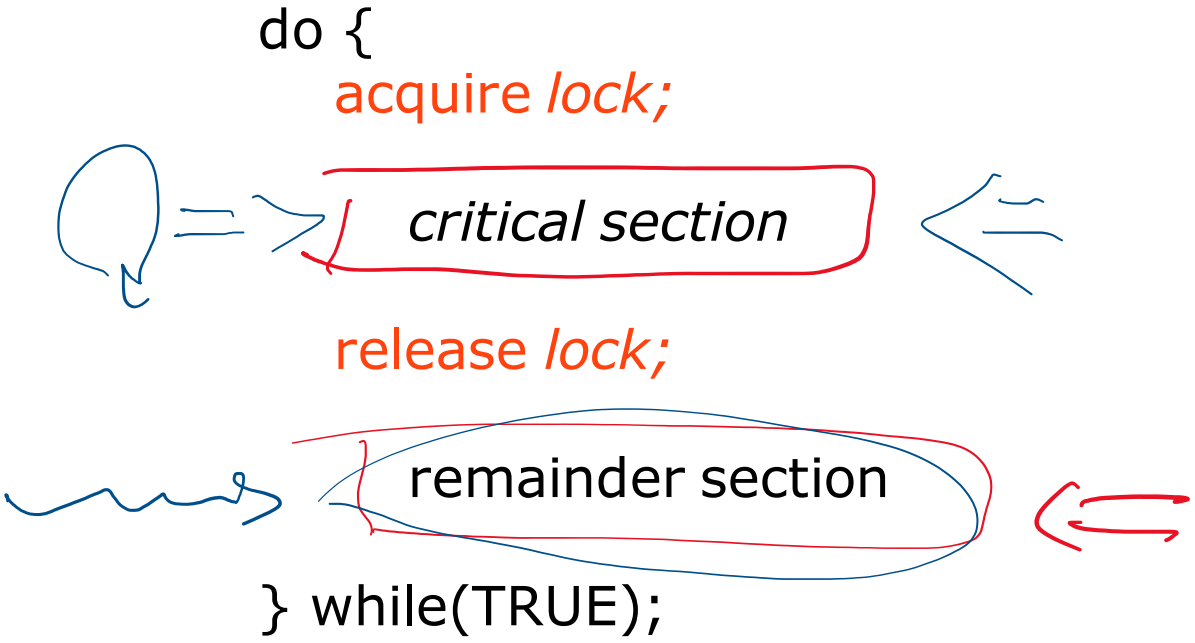

Scheduling and Mutual Exclusion

What happens when tasks share resources and use mutual exclusion to guard access to those resources?

Priority Inversion

- Priority based schedulers always run the task with higher priority

- What if a higher priority task is blocked by a lower priority task?

Priority Inversion: a scheduling anomaly where a high-priority task is blocked while unrelated lower-priority tasks are executing

Illustration of Priority Inversion

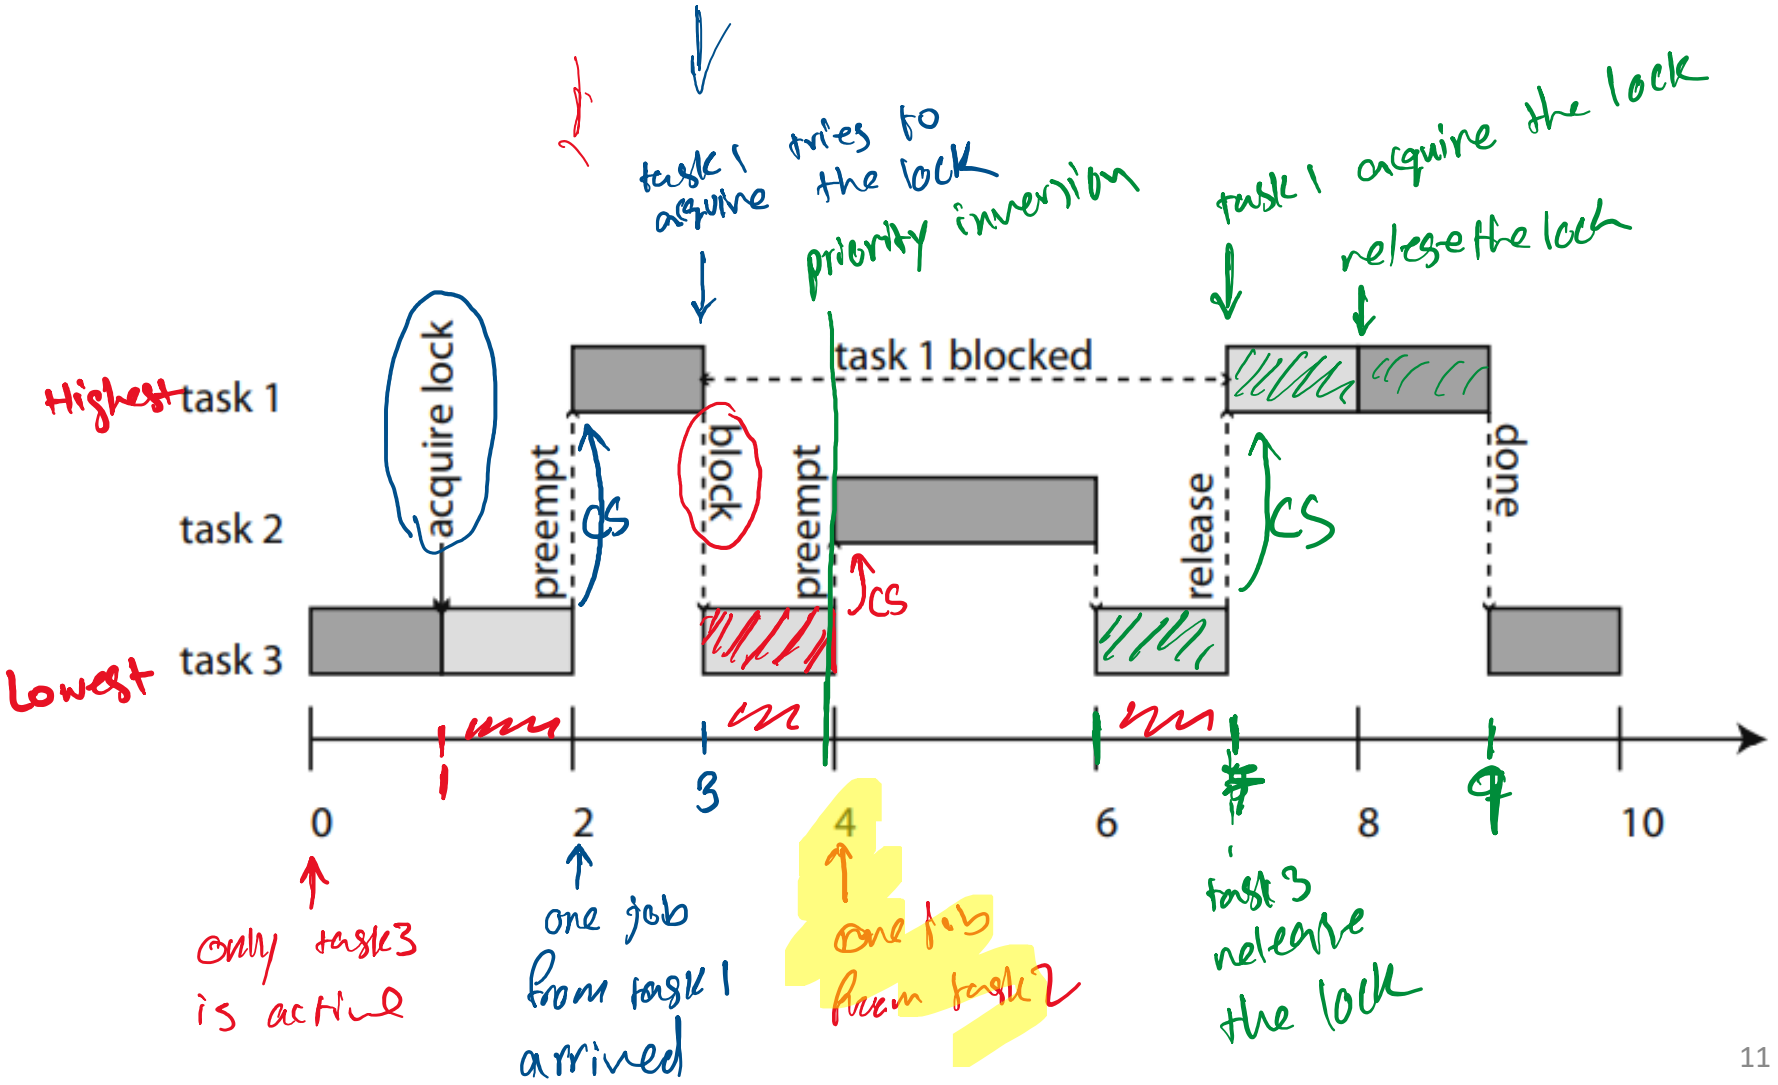

Priority Inheritance Protocol

- The task that holds the locks inherits the priority of the task that tries to acquire the lock

- Implication: prevents priority inversion
  - The task that holds the lock cannot be preempted by an intermediate priority task

Illustration of Priority Inheritances

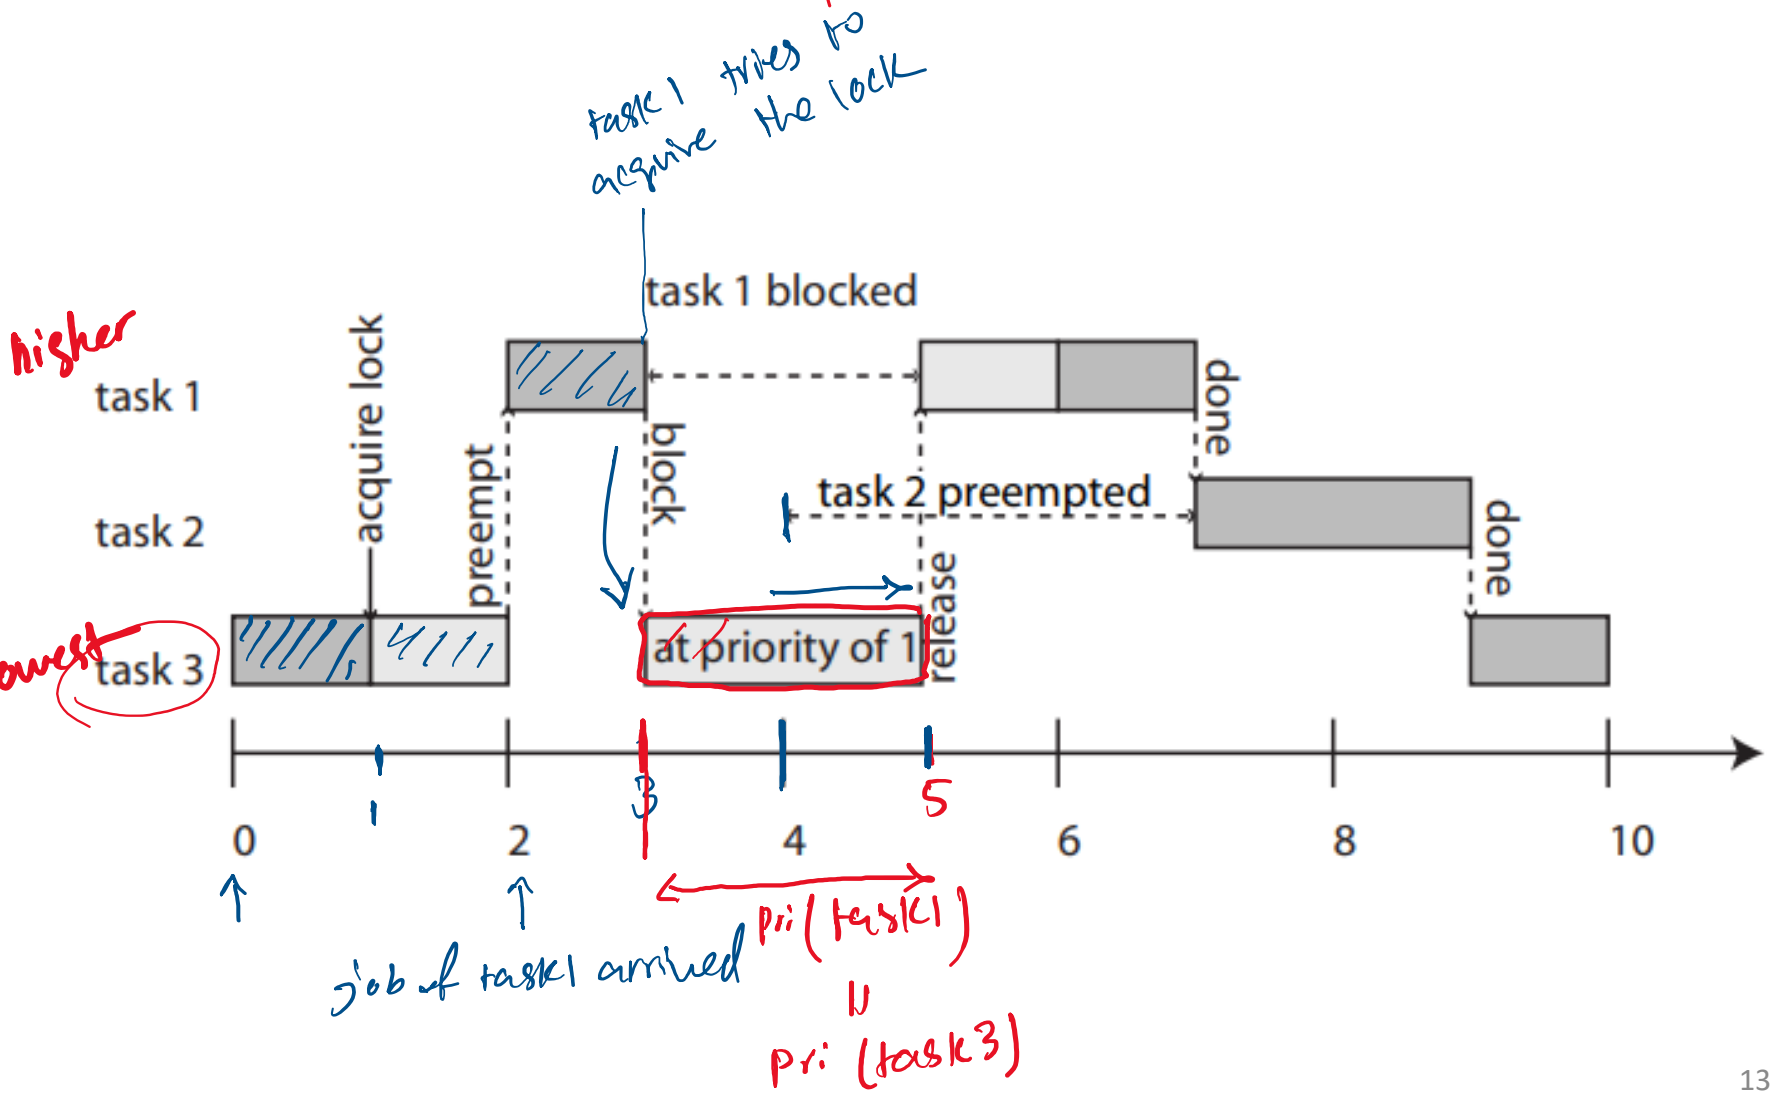

Deadlock!

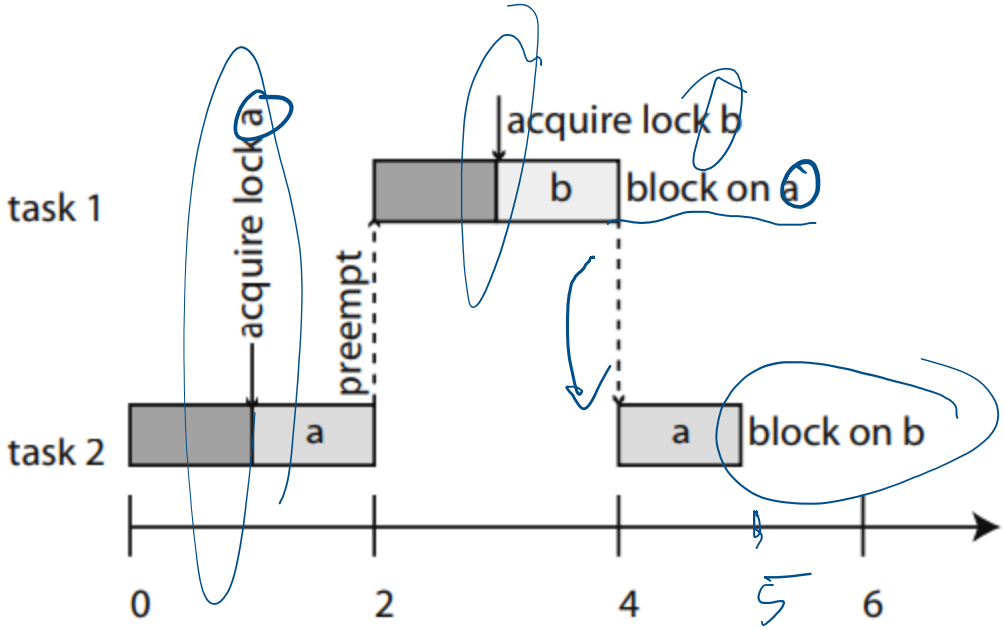
task 1
task 2
acquire lock a
acquire lock b
b
block on a
preempt
a
a
block on b
0
2
4
6

Priority Ceiling Protocol

- A task can acquire a lock only if the task's priority is strictly higher than the priority ceilings of all locks currently held by other tasks.
- PC(R) = highest priority of all tasks that may lock R
- Process P can start a new critical section only if: Task's priority > PC of all resources locked by other processes
- Assumption: we know the number of tasks and priorities

which task uses which lock

Illustration of Priority Ceiling Protocol

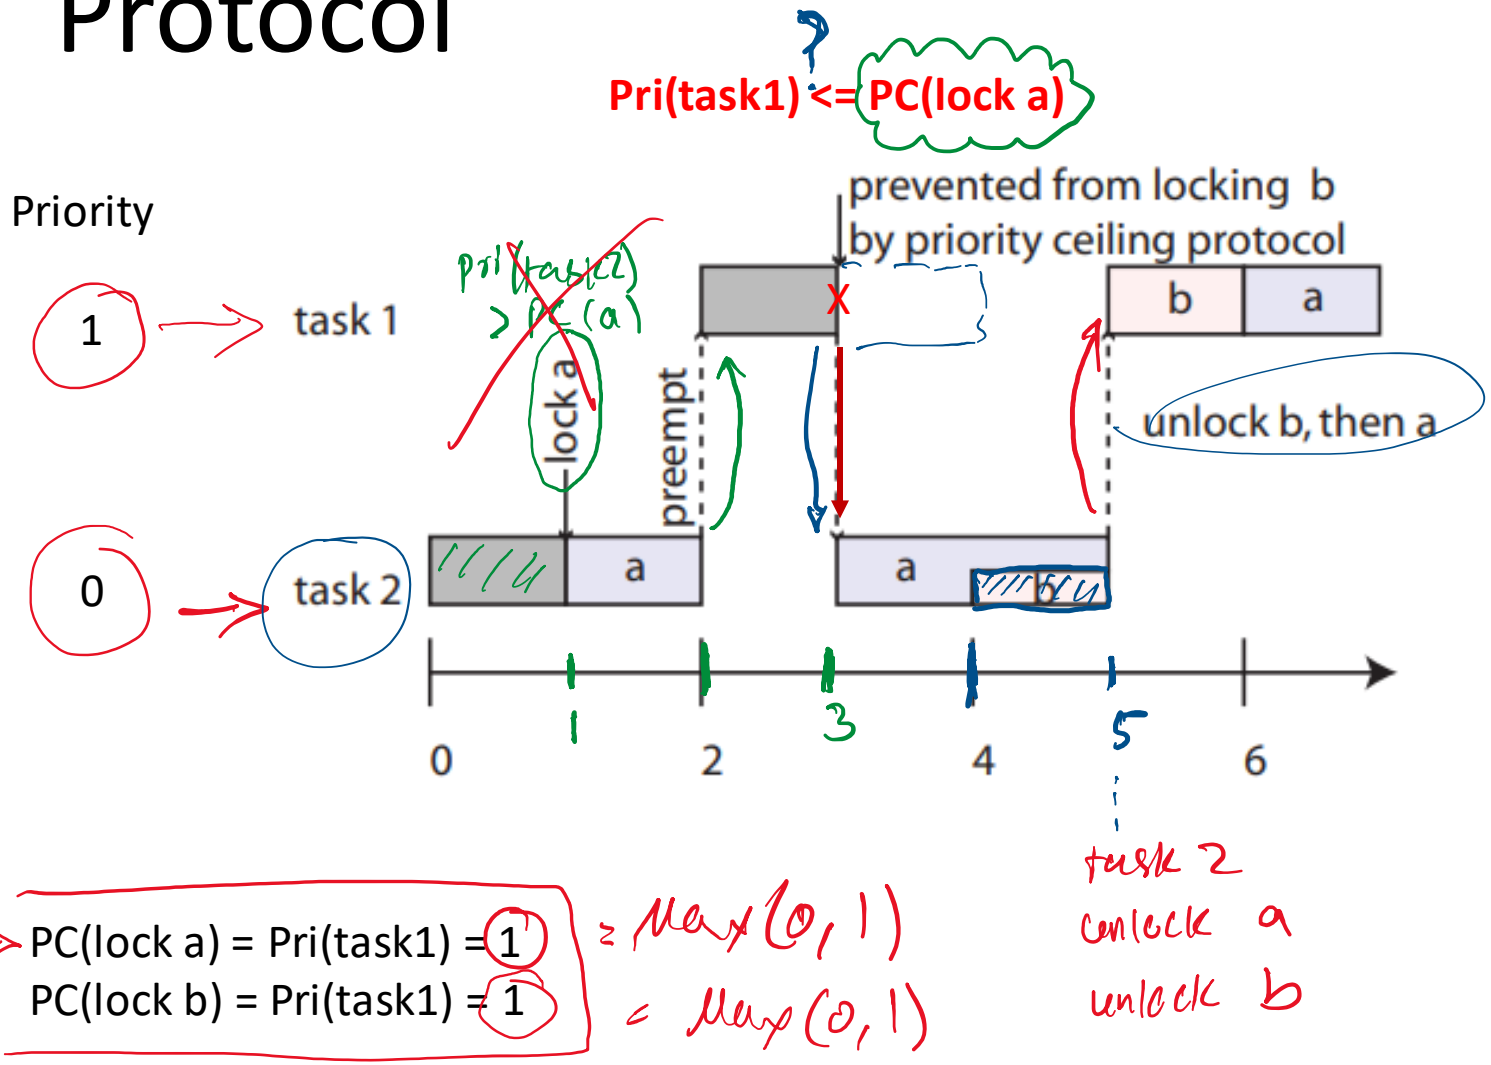

Another Example

PC(S1) = Pri(t2) = 2
PC(S2) = Pri(t3) = 1

$= \max(pri(\tau_2), pri(\tau_4)) = 2$ (2, 0)

$= \max(pri(\tau_3))$

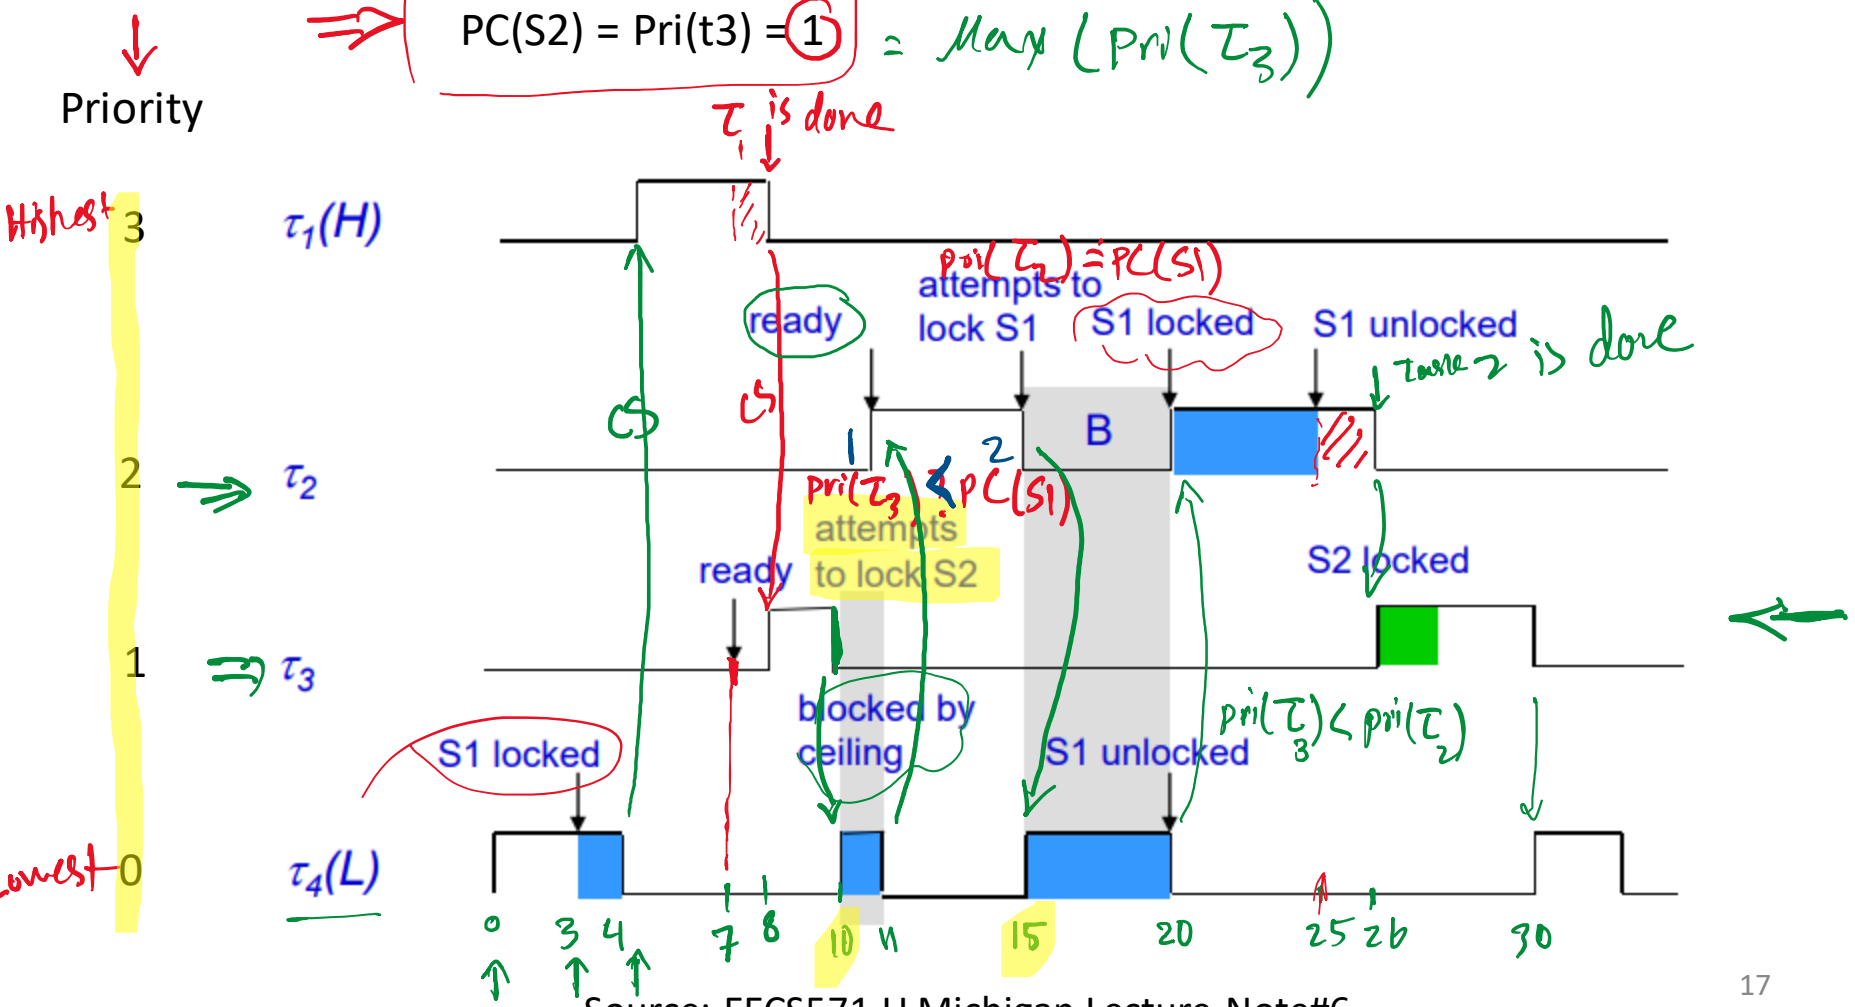

Source: EECS571 U Michigan Lecture Note#6

References

- Lee and Seshia. *Intro to Embedded Systems A Cyber Physical System Approach*., 2017. (Chapter 12)